

Learning Prostate Anatomy at Test Time for Cancer Detection in Micro-Ultrasound

Obed Korshie Dzikunu^{1,2}

OBED@ECE.UBC.CA

Mohammad Mahdi Abootorabi^{1,2}

MAHDI.ABOOTORABI@ECE.UBC.CA

Mohamed Harmanani^{2,3}

MOHAMED.HARMANANI@QUEENSU.CA

Paul F. R. Wilson^{2,3}

PAUL.WILSON@QUEENSU.CA

Emma Willis^{2,3}

EMMA.WILLIS@QUEENSU.CA

Ferdinand Luger⁴

FERDINAND.LUGER@ORDENSKLINIKUM.AT

Adam Kinnaird⁵

ASK@UALBERTA.CA

Brian Wodlinger⁶

BWODLINGER@EXACTIMAGING.COM

Parvin Mousavi^{2,3}

MOUSAVI@QUEENSU.CA

Purang Abolmaesumi¹

PURANG@ECE.UBC.CA

¹ Department of Electrical and Computer Engineering, University of British Columbia, Vancouver, Canada

² Vector Institute, Toronto, Canada

³ School of Computing, Queen's University, Kingston, Canada

⁴ Ordensklinikum Linz, Linz, Austria

⁵ Department of Surgery, University of Alberta, Edmonton, Canada

⁶ Exact Imaging, Markham, Canada

Abstract

Domain shift across clinical centers using different imaging hardware or acquisition protocols remains a fundamental barrier to deploying deep learning models for prostate cancer (PCa) detection. Existing test-time adaptation (TTA) methods address distribution shift through entropy minimization or augmentation-based self-supervision, correcting for statistical differences in image appearance but ignoring the anatomical structure of the target domain. We propose **ANT**, a segmentation-guided TTA framework that adapts a pretrained cancer detection encoder to the target domain by solving an auxiliary prostate segmentation task at test time, supervised by pseudo-masks from a frozen pretrained segmentation network. By aligning encoder representations to prostate anatomy in the target domain, ANT corrects domain-specific feature drift while preserving cancer-discriminative structure. The model was trained on 693 patients imaged with an earlier-generation micro-ultrasound scanner in a multi-center clinical trial, and evaluated on 118 patients acquired with a newer-generation system across two centers in another clinical trial. Under a leave-one-center-out protocol with identical evaluation conditions across all methods, ANT improves mean AUC by 2.9% and 3.6% at the biopsy-core and patient levels, respectively, over no adaptation, outperforming TTA baselines. Code is available at: <https://github.com/ObedDzik/ant.git>.

1. Introduction

Prostate cancer (PCa) remains one of the most prevalent malignancies among men worldwide (Rawla, 2019), and its diagnosis relies on histopathological analysis of tissue obtained via transrectal ultrasound (TRUS)-guided biopsy (Panzone et al., 2022). A fundamental challenge of conventional TRUS is its poor soft tissue contrast at standard frequencies,

which limits the ability of clinicians to visually localize suspicious regions of interest. As a result, systematic biopsy schemes sample multiple cores across the gland in a spatially uniform pattern rather than targeting suspected lesions, leading to frequent missed cancers and limiting diagnostic sensitivity (Mate et al., 2023). Pre-procedural multi-parametric MRI (mpMRI) partially addresses this by providing anatomical and functional information for targeted biopsy guidance and represents the current clinical gold standard for lesion localization (Ahmed et al., 2017). However, its high cost, logistical burden, and restricted availability limit its widespread adoption.

Micro-ultrasound addresses this gap by operating at higher transducer frequencies (approximately 29 MHz compared to conventional 5 – 9 MHz TRUS), producing spatial resolution sufficient for the detection of cancer similar to that of mpMRI at a fraction of the cost, while retaining the accessibility, portability, and real-time guidance capabilities of ultrasound (Rohrbach et al., 2018). Specifically, clinical studies have demonstrated that micro-ultrasound achieves sensitivity for clinically significant PCa detection comparable to mpMRI-guided biopsy (Eure et al., 2018; Klotz et al., 2020; Kinnaird et al., 2025), positioning it as a viable and scalable alternative for targeted biopsy guidance.

Deep learning models trained on micro-ultrasound have demonstrated the potential to automate cancer identification (Gilany et al., 2023; Wilson et al., 2024; Pensa et al., 2024; Willis et al., 2026) but models trained on data from a single clinical center may suffer significant performance degradation when deployed across different centers or scanner generations, a phenomenon known as domain shift (Quiñonero-Candela et al., 2009; Pooch et al., 2020). Test-time adaptation (TTA) offers a promising paradigm for addressing this challenge without requiring labeled target domain data. By adapting model parameters at inference time using a self-supervised auxiliary signal, TTA methods reduce the gap between training and deployment distributions (Sun et al., 2020; Wang et al., 2020). Existing TTA methods including TENT (Wang et al., 2020), SAR (Niu et al., 2023), EATA (Niu et al., 2022), and CoTTA (Wang et al., 2022), rely on entropy minimization or augmentation consistency as self-supervised signals. While effective for natural image distribution shifts, these signals do not exploit the structural and anatomical properties of medical images, properties which are clinically meaningful and tend to remain identifiable across domains even as image statistics vary substantially, as evidenced by the cross-domain success of generalist segmentation models (Ma et al., 2024; Wasserthal et al., 2023).

In prostate micro-ultrasound imaging, we hypothesize that the prostate gland remains a consistently identifiable anatomical target across patient populations, clinical centers, and scanner generations, even as low-level image statistics vary substantially due to differences in acquisition protocols and equipment. We further hypothesize that domain shift manifests as drift in the encoder’s internal feature representations of this anatomy, corrupting the features that underpin downstream cancer detection. We therefore propose that supervising the encoder to segment the prostate at test time re-anchors these representations to this domain-invariant anatomical target, correcting feature drift without requiring target domain labels. This anatomically grounded signal provides explicit structural supervision at test time, in contrast to generic self-supervised objectives that derive their adaptation signal from the model’s output distribution in the target domain, without leveraging domain-invariant properties of the input.

We therefore propose **ANT**, a TTA framework that adapts the encoder of a pretrained cancer detection model by solving a prostate segmentation auxiliary task at test time. A lightweight segmentation head, guided by pseudo-masks from a frozen pretrained segmentation network (MicroSegNet (Jiang et al., 2024)), generates a supervision signal that drives updates to the first n layers of a Vision Transformer (ViT) backbone. The segmentation head accumulates adaptation across the test set, while the encoder is restored to its checkpoint state after each sample to prevent cross-sample leakage. An overview of the proposed ANT framework during both the training and inference phases is illustrated in Figure 1.




Figure 1: Overview of ANT framework. At inference, the first n layers of the trained cancer detection encoder is domain-adapted via an auxiliary segmentation task guided by pseudo-masks from a frozen segmentation network. The adapted encoder is then used for cancer detection.

Our contributions are as follows:

1. We propose a segmentation-guided TTA framework that leverages anatomical invariance to improve prostate cancer detection in micro-ultrasound. To the best of our knowledge, this is the first work to utilize organ segmentation as a test-time adaptation signal for cancer detection.
2. Across two multi-center clinical trial cohorts, we show that the proposed method consistently outperforms existing TTA approaches and improves upon a prior state-of-the-art prostate cancer detection model.
3. Through controlled ablation studies, we show that anatomical segmentation provides a robust adaptation signal for mitigating domain shift arising from variations in imaging data distributions.

Generalizable Insights about Machine Learning in the Context of Healthcare

This work offers the following insights relevant to machine learning in healthcare beyond the specific application of prostate cancer detection:

- Anatomical invariance as a test time adaptation signal: In medical imaging, anatomical structures remain consistently identifiable across domains even when image statistics differ substantially, as evidenced by the success of generalist segmentation models that transfer across scanners, centers, and imaging modalities without domain-specific retraining (Ma et al., 2024; Wasserthal et al., 2023). The results of this work suggest that auxiliary tasks exploiting this identifiability, such as organ segmentation, can provide more targeted and stable adaptation signals than generic self-supervised objectives. While demonstrated here for prostate micro-ultrasound, we hypothesize that this principle may extend to other medical imaging applications where a domain-invariant anatomical supervision target can be defined.
- Domain shift is encoded in early encoder layers: Our ablation results highlight that updating the first n encoder layers, which process low-level signal statistics, outperforms updating middle or late layers. This suggests that cross-site and hardware-induced distribution shift may manifest as low-level feature perturbations, making the initial encoder layers an effective target for test-time adaptation.

2. Related Work

2.1. Learning Strategies for Domain Shift

A variety of approaches have been proposed to improve robustness under domain shift, broadly spanning domain adaptation, domain generalization, and meta-learning (Guan and Liu, 2021; Yoon et al., 2024). Domain adaptation methods assume access to labeled source data and unlabeled target data during training, and aim to reduce distribution mismatch via feature alignment, adversarial learning, or image translation (Ganin et al., 2016; Kamnitsas et al., 2017; Chen et al., 2020). Domain generalization methods on the other hand learn representations that generalize to unseen domains without access to target data, using strategies such as invariant feature learning (Chen et al., 2023) and data augmentation (Zhang et al., 2020; Zhou et al., 2024). Meta-learning approaches further simulate domain shift during training by optimizing across source domain splits to improve robustness to unseen domains (Finn et al., 2017; Li et al., 2018). While effective, these methods rely on training-time assumptions about domain variability and may not generalize to deployment settings where domain characteristics are unknown or evolving.

In prostate micro-ultrasound, domain shift may arise from differences in scanner generation, transducer frequency, and center-specific acquisition protocols, which may degrade model performance under deployment conditions. Existing approaches primarily address this through training-time design. For example, ProstNFound (Wilson et al., 2024) integrates foundation models with ultrasound-specific features and clinical biomarkers to achieve strong multi-center performance on prostate cancer detection, while ProstNFound+ (Wilson et al., 2026) demonstrates improved generalisation to data acquired years later at a new clinical site. Although these methods achieve strong performance across the evaluated

cohorts, they remain fixed after training and are therefore limited to the domain variability observed during development. Similarly, while large-scale foundation models have shown robustness to cross-center variation in prostate MRI (Li et al., 2024), their behaviour under unseen micro-ultrasound acquisition conditions remain less well characterized. This motivates test-time adaptation methods, such as DEnEM (Gilany et al., 2024), which adapt model predictions at inference time using only incoming data to better handle previously unseen domain shifts.

2.2. Test-Time Adaptation

TTA adjusts models at inference using unlabeled target-domain data, without access to source labels or retraining (Liang et al., 2024). TENT (Wang et al., 2020) established the entropy minimization paradigm by updating batch normalization parameters at test time. Building on this, EATA (Niu et al., 2022) introduces sample selection and Fisher regularization to prevent catastrophic forgetting, while SAR (Niu et al., 2023) improves stability under small batches via sharpness-aware minimization. CoTTA (Wang et al., 2022) extends TTA to continually shifting distributions through stochastic weight restoration and augmentation-averaged predictions, and ROID (Marsden et al., 2024) further addresses universal TTA through diversity-weighted loss and continual weight ensembling. MEMO (Zhang et al., 2022) takes a different approach, adapting each test sample independently via marginal entropy minimization over augmented views.

A recent large-scale benchmark of TTA methods in medical image segmentation (Yu et al., 2025) evaluated twenty representative methods across seven imaging modalities under unified protocols, finding that no single paradigm generalizes across all conditions and that entropy-based methods degrade significantly under large inter-center and inter-device shifts. This finding is consistent with our observations in multi-center micro-ultrasound and motivates the need for domain-specific adaptation signals beyond generic entropy objectives.

More recently, task-specific TTA methods have been proposed for medical imaging. GraTa (Chen et al., 2025) improves adaptation by aligning pseudo-task gradients with the primary task gradient direction using dynamic learning rates. A3-TTA (Wu et al., 2025) applies anchor-guided pseudo-label generation with boundary-aware entropy minimization for prostate segmentation under domain shift. DEnEM (Gilany et al., 2024) addresses domain shift in micro-ultrasound prostate cancer detection through neural network ensembles for calibrated marginal entropy minimization. ReservoirTTA (Vray et al., 2025) maintains domain-specialised model reservoirs with online style clustering for recurring distribution shifts.

Despite this progress, existing TTA methods rely on entropy minimization, augmentation consistency, or gradient alignment as adaptation signals, all of which operate on the global feature statistics or output logits without exploiting domain-invariant structural properties of the input. We address this gap by using prostate segmentation as a structured, anatomically grounded TTA signal that directly targets domain-specific feature drift in the encoder, providing a principled alternative to entropy-based adaptation for clinical deployment under multi-center domain shift.

3. Methods

3.1. Problem Formulation

Let $f_{\text{enc}} : \mathcal{X} \rightarrow \mathcal{Z}$ denote a pretrained encoder mapping micro-ultrasound images to feature representations, and let $g : \mathcal{Z} \rightarrow \mathcal{Y}$ denote the detection head mapping features to cancer probability heatmaps. The model $h = g \circ f_{\text{enc}}$ is trained on source domain data $\mathcal{D}_{\text{src}} = \{(x_i, y_i)\}_{i=1}^N$ from a clinical trial cohort noted as Center A.

At test time, the model encounters target domain data $\mathcal{D}_{\text{tgt}} = \{x_j\}_{j=1}^M$ from unseen centers in a different cohort with no available labels. The domain gap $\mathcal{G} = d(\mathcal{D}_{\text{src}}, \mathcal{D}_{\text{tgt}})$ degrades detection performance. Our goal is to adapt f_{enc} to $\mathcal{D}_{\text{tgt}}$ using only unlabeled test samples and an anatomical auxiliary task.

3.2. Anatomical Alignment via Segmentation TTA

Pseudo-label generation. We use a pretrained frozen MicroSegNet (Jiang et al., 2024) $s_\phi : \mathcal{X} \rightarrow [0, 1]^{H \times W}$ to generate prostate pseudo-masks for each test image:

$$\hat{m}_j = \sigma(s_\phi(x_j)), \quad (1)$$

where σ denotes the sigmoid function. The MicroSegNet parameters ϕ are frozen throughout TTA.

Segmentation head. A lightweight segmentation head $p_\psi : \mathcal{Z} \rightarrow \mathbb{R}^{H' \times W'}$ is attached to the encoder output:

$$\hat{s}_j = p_\psi(f_{\text{enc}}(x_j)), \quad (2)$$

where $\hat{s}_j$ is upsampled to match $\hat{m}_j$ via bilinear interpolation. The head consists of two convolutional layers with GroupNorm and a final 1×1 projection to a single channel binary output. The head is randomly initialised and is not reset between test samples, allowing it to accumulate an increasingly refined understanding of target domain anatomy across the test set.

TTA objective. The TTA loss combines Dice loss and binary cross-entropy:

$$\mathcal{L}_{\text{TTA}} = \mathcal{L}_{\text{Dice}}(\hat{s}_j, \hat{m}_j) + \lambda \mathcal{L}_{\text{BCE}}(\hat{s}_j, \hat{m}_j), \quad (3)$$

where λ balances the two terms. Dice loss encourages global shape agreement between predicted and pseudo segmentation masks, while binary cross-entropy enforces pixel-level alignment.

3.3. Encoder Update Strategy

Rather than updating all encoder parameters, we selectively update all parameters of the first n transformer blocks of the ViT backbone:

$$\Theta_{\text{update}} = \{\theta_l : l \in \mathcal{I}_f\}, \quad (4)$$

where θ_l denotes all parameters of block l and $\mathcal{I}_f$ denotes the indices of the first n blocks.

Algorithm 1: ANT: Segmentation-Guided Test-Time Adaptation

Input: Image x_j , checkpoint W^* , frozen segmentor s_ϕ , steps τ , lr η
Persistent: Seg head ψ , optimizer $\mathcal{O}$
Output: Cancer score $\hat{y}_j$

```

 $\hat{m}_j \leftarrow \sigma(s_\phi(x_j))$  // pseudo-mask
 $W_{\text{enc}} \leftarrow W^*$  // initialise from checkpoint
for  $t \leftarrow 1$  to  $\tau$  do // adapt first  $n$  blocks
     $\hat{s}_j \leftarrow \text{Up}(p_\psi(f_{W_{\text{enc}}}(x_j)))$ 
     $\mathcal{L} \leftarrow \mathcal{L}_{\text{Dice}}(\hat{s}_j, \hat{m}_j) + \lambda \mathcal{L}_{\text{BCE}}(\hat{s}_j, \hat{m}_j)$ 
     $W_{\text{enc}}, \psi \leftarrow \mathcal{O}(\nabla_{\Theta, \psi} \mathcal{L})$ 
end

 $\hat{y}_j \leftarrow g(f_{W_{\text{enc}}}(x_j))$  // predict with adapted encoder
 $W_{\text{enc}} \leftarrow W^*$  // reset encoder,  $\psi$  and  $\mathcal{O}$  persist
return  $\hat{y}_j$ 

```

3.4. PCa Detection

3.4.1. DETECTION MODEL ARCHITECTURE

The cancer detection model consists of three components. The image encoder is a pre-trained DINOv3 ViT-L/16 backbone (Siméoni et al., 2025) with 24 transformer blocks, each equipped with LayerNorm layers. Intermediate hidden states are extracted at interaction indices $\{4, 11, 17, 23\}$ and processed by a UNETR-style hierarchical decoder (Hatamizadeh et al., 2022) comprising skip-connection encoder and decoder blocks that progressively up-sample features to produce a spatial cancer probability heatmap.

The heatmap features are subsequently passed to a two-way transformer (Kirillov et al., 2023) that attends jointly over image embeddings and learnable query tokens, enabling bidirectional cross-attention between spatial features and task-specific representations. The output is processed by two task-specific decoders: a mask decoder that produces a dense cancer probability heatmap over the biopsy region, and a class decoder that produces an image-level classification score for clinically significant PCa. Cancer predictions at the biopsy core level are obtained by averaging heatmap values along the biopsy needle trajectory.

During TTA, only the first n transformer blocks of the DINOv3 encoder are updated via the segmentation auxiliary task. The UNETR decoder, two-way transformer, mask decoder, and class decoder remain frozen throughout adaptation.

3.4.2. TRAINING OBJECTIVE

The detection model is trained using the multi-task objective proposed by Wilson et al. (2026), combining image-level classification supervision with weak spatial constraints on the predicted heatmap. Image-level supervision is applied to the csPCa probability via a binary cross-entropy loss with the pathology label. In parallel, spatial supervision is derived from core-level involvement labels by aggregating the predicted heatmap within the biopsy

needle region and comparing it to the corresponding involvement score using binary cross-entropy. The final objective is the sum of these two terms, jointly encouraging accurate core-level classification while promoting spatially coherent heatmap responses that reflect tumour burden within the sampled region.

Let $y_{\text{cls}} \in \{0, 1\}$ denote the core-level pathology label and $\hat{y}_{\text{cls}} \in [0, 1]$ the predicted probability. The model outputs a spatial heatmap $P \in [0, 1]^{H \times W}$, which is aggregated over the biopsy region defined by a binary mask $M \in \{0, 1\}^{H \times W}$. Let $\mathcal{N} = \{(i, j) \mid M_{i,j} = 1\}$ denote the set of valid pixels within the needle region. The predicted cancer involvement is computed as a masked expectation:

$$\hat{I} = \frac{1}{|\mathcal{N}|} \sum_{(i,j) \in \mathcal{N}} P_{i,j}. \quad (5)$$

The training objective combines core-level and region-level supervision:

$$\mathcal{L} = \mathcal{L}_{\text{csPCa}} + \mathcal{L}_{\text{map}}, \quad (6)$$

where

$$\mathcal{L}_{\text{csPCa}} = \text{BCE}(y_{\text{cls}}, \hat{y}_{\text{cls}}), \quad \mathcal{L}_{\text{map}} = \text{BCE}(I, \hat{I}), \quad (7)$$

and $I \in [0, 1]$ denotes the ground-truth cancer involvement score.

4. Study Cohort

4.1. Cohort Selection

Data were derived from clinical trials with predefined inclusion and exclusion criteria. Enrolled patients were indicated for prostate biopsy due to elevated PSA and/or abnormal digital rectal examination, were aged ≥ 18 years, and had no prior prostate biopsy or history of genitourinary cancer. Additional criteria included no contraindications to biopsy or mpMRI, and no recent mpMRI for prostate cancer evaluation. Exclusion criteria mirrored these conditions.

4.2. Data Acquisition and Study Design

The training cohort consisted of retrospective data collected between 2013 and 2016 from a multi-center clinical trial (NCT02079025) involving five institutions, collectively called Center A (source). Patients underwent untargeted systematic prostate biopsy using an earlier version of the ExactVu μ US system (Exact Imaging, Markham, Canada), with typically 10 – 12 cores acquired per subject (Rohrbach et al., 2018).

The test cohort comprised prospective data collected between 2021 and 2024 (NCT05220501) from two clinical sites using a newer version of the ExactVu μ US system. In this cohort, patients underwent both systematic and targeted biopsies. Targeting was performed using either micro-ultrasound (μ US) with PRI-MUS scoring or pre-procedural mpMRI with PI-RADS scoring, depending on the study arm. MRI- μ US fusion was conducted using the ExactVu FusionVu system (Kinnaird et al., 2025).

Biopsy locations were assigned PRI-MUS-based risk scores across both arms. In the MRI-guided arm, targeted cores inherited the PI-RADS score of the corresponding lesion;

Table 1: Cohort statistics per clinical center. Centers differ in scanner generation and acquisition protocol, representing realistic deployment-time domain shift.

Center	Patients	Cores	Cancer %
A (source)	693	5479	16%
B	76	1118	34.6%
C	42	569	27.0%

if no lesion was identified, a prostate-level PI-RADS score (1 or 2) from the radiology report was assigned to all cores from that patient. Summary statistics for both cohorts are provided in Table 1.

4.3. Data Extraction

Sagittal B-mode micro-ultrasound images were acquired at 28 mm depth and 46.06 mm width during ultrasound-guided biopsy. A 30-second cineloop was recorded per core, capturing needle alignment and firing. The needle trace was defined by annotating the tip position and insertion angle at the moment of firing, from which a corresponding path mask was generated. The final frame immediately prior to needle entry, co-registered with the needle trace, was selected as the representative image for each biopsy core. This frame captures the pre-fire tissue state without needle artefact and is the standard input for learning-based analysis in this setting (Wilson et al., 2026).

Histopathological analysis of each biopsy specimen provided the cancer diagnosis, percentage involvement (approximate fraction of tumour length within the core), and ISUP Grade Group (GG 0 – 5). Clinically significant PCa (csPCa) was defined as $\text{GG} \geq 3$; cores with GG 1 – 2 were categorised as insignificant PCa and GG 0 as benign. Labels were assigned to the full needle trace region, introducing a minor spatial approximation for partially involved cores. Patient-level labels reflected the most severe core finding for that patient.

4.4. Feature Choices

Raw B-mode images were resized from their native resolution of 1372×833 pixels to 512×512 pixels using bilinear interpolation and normalised to the range $[0, 1]$. All predictions were made solely from imaging features with no clinical metadata provided as input.

The resized images were passed to the DINOv3 ViT-L/16 encoder, which divides each image into non-overlapping patches of 16×16 pixels, yielding a sequence of $32 \times 32 = 1024$ patch tokens each of dimension 1024. These patch tokens are processed by the transformer blocks to produce contextualized spatial feature representations, which are extracted at interaction indices $\{4, 11, 17, 23\}$ and passed to the UNETR decoder for heatmap generation. Cancer predictions at the biopsy core level are obtained by averaging heatmap activations along the needle trace mask $\mathcal{N}$ as described in Section 3.4.2.

Table 2: AUC (95% bootstrap CI) for all methods across both evaluation centres (B and C). Bold indicates the best performance per column.

Method	Patient-Level csPCa AUC		Core-Level csPCa AUC		Core-Level PCa AUC	
	B	C	B	C	B	C
No adaptation	77.6(72.9 – 82.0)	79.4(66.8 – 89.9)	78.5(73.6 – 83.3)	84.5(74.0 – 93.1)	71.5(68.2 – 74.6)	73.8 (68.5 – 79.3)
MEMO (Zhang et al., 2022)	78.3(73.8 – 82.5)	83.5(73.0 – 91.7)	77.7(73.1 – 82.5)	83.4(71.8 – 92.9)	73.3(70.2 – 76.3)	70.6(65.3 – 76.3)
TENT (Wang et al., 2020)	76.0(71.6 – 80.2)	75.0(60.8 – 87.0)	80.1(76.0 – 84.3)	81.9(70.3 – 92.6)	74.5(71.5 – 77.3)	68.6(62.7 – 74.4)
EATA (Niu et al., 2022)	73.5(68.7 – 77.2)	69.4(65.1 – 72.8)	79.8(75.3 – 82.4)	80.5(77.3 – 84.2)	74.9 (71.0 – 77.2)	71.9(69.8 – 74.5)
CoTTA (Wang et al., 2022)	77.9(73.2 – 82.2)	77.9(64.6 – 89.7)	78.5(73.7 – 83.3)	84.1(73.8 – 92.7)	71.4(68.2 – 74.5)	72.9(67.4 – 78.4)
ROID (Marsden et al., 2024)	78.2(73.8 – 82.2)	76.8(64.1 – 87.6)	79.9(75.4 – 84.3)	84.4(73.0 – 92.9)	74.3(71.3 – 77.3)	71.0(65.3 – 76.6)
SAR (Niu et al., 2023)	79.5(75.2 – 83.7)	79.2(68.1 – 89.5)	80.9(76.4 – 85.5)	82.5(72.5 – 91.1)	73.3(70.3 – 76.4)	70.5(64.9 – 75.8)
ANT (ours)	79.8 (75.6 – 83.7)	84.3 (73.4 – 92.8)	81.0 (76.8 – 85.0)	*87.8 (76.4 – 96.0)	73.0(69.9 – 76.2)	72.3(66.7 – 77.7)

*Core-Level csPCa AUC at Centre C: ANT vs SAR, DeLong test $p = 0.044$.

5. Results

5.1. Evaluation Approach and Study Design

The detection model is trained exclusively on Center A and evaluated under a leave-one-center-out (LOCO) protocol over target centers B and C, with one serving as the validation set for checkpoint selection and the other as the test set. Performance is reported as biopsy-core-level area under the curve value (AUC). At no point are target labels used to supervise adaptation; all methods operate exclusively on unlabeled test images. All baselines are reported using their published hyperparameter configurations. However, for EATA(Niu et al., 2022) and SAR(Niu et al., 2023), the entropy margin was rescaled from the ImageNet-calibrated value to account for binary classification following the same proportional logic as the original papers. TENT(Wang et al., 2020), CoTTA(Wang et al., 2022), ROID(Marsden et al., 2024), and MEMO(Zhang et al., 2022) required no such rescaling and are reported directly under their published defaults.

5.2. Comparison Against TTA Baselines

Table 2 reports AUC on held-out test centers under the LOCO protocol. ANT achieves the highest mean core-level AUC and patient-level AUC for clinically significant cancer of 84.4% and 82.1%, respectively, outperforming the no-adaptation baseline by 2.9% and 3.6% at the core and patient level, and surpassing all competing TTA methods. Notably, ANT achieves particularly strong performance on Center C, reaching 87.8% core-level AUC and 84.3% patient-level AUC and a sensitivity of 94.4% and 83.1% at specificity values of 60% and 80%, respectively, at the core level (Table 3).

Existing TTA methods show mixed and inconsistent results relative to the no-adaptation baseline. All baseline methods tend to improve on one center while degrading on the other, resulting in mean AUCs that are comparable to or below the no-adaptation baseline across most metrics. For instance, TENT achieves 80.1% core-level csPCa AUC on Center B but drops to 81.9% on Center C, and SAR shows a similar pattern with 80.9% on Center B falling to 82.5% on Center C. Moreover, ANT achieves statistically significant improvement over SAR, the second-best baseline, on core-level csPCa AUC at Center C. We attribute the superiority of ANT to the anatomical alignment signal, which provides a structured and domain-invariant supervision target that generic self-supervised signals cannot replicate.

Table 3: Sensitivity at fixed specificity (95% bootstrap CI) for all methods across both evaluation centres (B and C) at the core level for clinically significant prostate cancer (csPCa). Bold indicates the best performance per column.

Method	Sensitivity @ 60% Specificity (%)		Sensitivity @ 80% Specificity (%)	
	B	C	B	C
No adaptation	81.0(73.5 – 88.2)	85.3(43.7 – 100.0)	67.1(56.1 – 76.5)	70.1(43.7 – 94.1)
MEMO (Zhang et al., 2022)	82.5(75.3 – 89.4)	85.2(64.7 – 100.0)	61.7(51.3 – 72.1)	74.5(50.0 – 95.6)
TENT (Wang et al., 2020)	80.4(72.7 – 87.5)	83.1(47.1 – 100.0)	66.4(57.8 – 75.4)	71.1(47.1 – 91.7)
EATA (Niu et al., 2022)	80.2(71.3 – 86.4)	89.5(49.2 – 100.0)	65.5(57.3 – 73.2)	63.2(44.6 – 88.5)
CoTTA (Wang et al., 2022)	80.0(72.0 – 87.3)	83.4(53.3 – 100.0)	65.8(55.8 – 75.5)	75.2(53.3 – 94.4)
ROID (Marsden et al., 2024)	83.4 (76.7 – 90.2)	94.0(55.6 – 100.0)	68.7(60.2 – 77.2)	77.4(55.6 – 95.9)
SAR (Niu et al., 2023)	81.9(74.5 – 88.8)	88.5(43.5 – 100.0)	70.9 (62.5 – 79.5)	72.2(43.5 – 94.1)
ANT (ours)	83.4 (76.8 – 90.1)	94.4 (64.3 – 100.0)	69.7(60.5 – 78.5)	83.1 (64.3 – 100.0)

Table 4: Per-patient AUC for ANT and the unadapted baseline stratified by MicroSegNet pseudo-mask Dice score tertile. Values reported as median (range). csPCa AUC is available only for patients with at least one positive and one negative core for clinically significant cancer.

Dice Tertile	n	Dice Median (Range)	PCa AUC (%)			csPCa AUC (%)		
			ANT	Baseline	ANT > Base	ANT	Baseline	ANT > Base
Low (≤ 0.72)	11	0.60(0.36 – 0.71)	80.0(29.6 – 100.0)	78.6(23.3 – 100.0)	4/11	100.0(78.1 – 100.0)	80.0(65.6 – 82.7)	3/3
Mid (0.72 – 0.90)	10	0.83(0.72 – 0.90)	68.3(6.2 – 100.0)	76.2(18.2 – 100.0)	3/10	63.6(50.0 – 95.8)	54.5(36.7 – 70.8)	5/5
High (> 0.90)	11	0.95(0.91 – 0.99)	57.1(21.4 – 94.4)	44.0(21.4 – 96.3)	6/11	71.7(40.0 – 100.0)	65.7(23.1 – 100.0)	4/6
Overall	32	0.83(0.36 – 0.99)	-	-	13/32	-	-	12/14

5.3. Sensitivity to Pseudo-Mask Quality

To assess ANT’s sensitivity to pseudo-mask quality, ground truth prostate segmentation masks were obtained for 32 randomly sampled cases spanning both evaluation centres. MicroSegNet Dice scores were computed against these masks and ranged from 0.36 to 0.99 (median 0.83). Cases were stratified into tertiles based on Dice score: low (≤ 0.72), mid (0.72 – 0.90), and high (> 0.90). Per-patient AUC was computed for PCa and csPCa separately for ANT and the unadapted baseline.

Table 4 summarizes the results. Despite substantially imperfect segmentation in the low Dice tertile (median Dice = 0.60, range 0.36 – 0.71), ANT consistently improves over the unadapted baseline on the clinically primary metric, csPCa AUC, across all three tertiles (3/3, 5/5, and 4/6 cases respectively; 12/14 overall). Of the 16 cases (similar performance on 3 cases) where ANT underperforms on PCa AUC, ANT recovers superiority on csPCa AUC in 5/6 of those cases with available csPCa data. Thus, ANT’s adaptation signal remains beneficial for clinically significant disease detection across the full range of observed mask quality.

5.4. Generalization to State-of-the-Art Architecture

To demonstrate that ANT is model-independent, we apply it to ProstNFound+ (PNF+) (Wilson et al., 2026), the current state-of-the-art model for prostate cancer detection in

Table 5: ANT applied to ProstNFound+ (PNF+), the current state-of-the-art model for micro-ultrasound prostate cancer detection. Results (AUC, %) demonstrate model-independent generalization of the proposed TTA approach. Best results in bold.

Method	Patient-Level csPCa AUC		Core-Level csPCa AUC		Core-Level PCa AUC	
	Center B	Center C	Center B	Center C	Center B	Center C
PNF+ (Wilson et al., 2026)	71.2(66.8 – 75.6)	75.6(64.4 – 85.3)	72.4(68.2 – 76.5)	75.4(62.8 – 87.2)	69.8(66.6 – 73.1)	63.8(58.1 – 69.4)
PNF+ + ANT (ours)	72.8 (68.3 – 77.4)	80.5 (69.7 – 88.8)	73.9 (69.8 – 78.0)	78.2 (66.5 – 87.1)	*73.1 (70.4 – 76.3)	65.2 (59.3 – 70.5)

*Core-Level PCa AUC at Centre B: ANT vs PNF+, DeLong test $p < 0.001$.

Table 6: Sensitivity at fixed specificity (95% bootstrap CI) for the unadapted ProstNFound baseline (PNF) and ANT across both evaluation centres (B and C) at the core level for clinically significant prostate cancer (csPCa).

Method	Sensitivity @ 60% Specificity (%)		Sensitivity @ 80% Specificity (%)	
	B	C	B	C
PNF+ (Wilson et al., 2026)	76.4(68.6 – 84.3)	72.3(50.0 – 92.9)	44.5(33.3 – 56.0)	58.4 (30.4 – 87.5)
PNF+ + ANT (ours)	78.9 (70.6 – 86.4)	85.5 (65.0 – 100.0)	46.0 (36.2 – 56.7)	56.0(30.8 – 77.8)

micro-ultrasound. PNF+ uses a MedSAM (Ma et al., 2024) image encoder with adapter-based fine-tuning and a prompt-driven class decoder. We apply ANT to the PNF+ encoder following the same protocol as in Section 3.3, updating the first n encoder blocks using segmentation pseudo-masks from the frozen MicroSegNet, without retuning any hyperparameters.

Tables 5 and 6 report performance of PNF+ with and without ANT on the held-out test centers. ANT improves PNF+ performance by a mean AUC of at least 3% across test centers, confirming that anatomical alignment via segmentation provides a consistent domain adaptation benefit regardless of the underlying detection architecture. Notably, this improvement is achieved without any architecture-specific tuning.

Figure 2 shows the AUROC for clinically significant cancer detection across increasing involvement thresholds and the distribution of activation leakage stratified by tissue type and model, for both Center B and Center C. Both models demonstrate improved discriminative performance as the involvement threshold increases, with AUROC rising from approximately 70 – 74% at threshold 0% to 87 – 88% at 70% for Center B, and from 63 – 67% to 79 – 80% for Center C. PNF+ + ANT consistently outperforms PNF+ across all thresholds and both centers, suggesting that the proposed test time adaptation method improves the model’s ability to distinguish clinically significant cancer from benign cores, particularly at lower involvement thresholds where the task is hardest. Regarding spatial specificity, activation leakage, defined as the ratio of average prostate heatmap activation to average needle-region activation, is markedly reduced by ANT across both centers and tissue types. While PNF+ exhibits a pronounced upper tail in the leakage distribution, reflecting frequent off-target activations outside the biopsy needle region, PNF+ + ANT concentrates nearly all of its mass near zero with no comparable tail, indicating substantially tighter spatial confinement of model activations. This improvement in anatomical speci-




Figure 2: Left: AUROC for clinically significant cancer prediction evaluated across increasing involvement thresholds (cores with true involvement $\geq t$ or benign). Right: Distribution of activation leakage, defined as the ratio between average prostate heatmap activation and average needle-region activation, stratified by cancer status (benign vs PCa) and method, PNF+ (no TTA) vs. PNF+ + ANT. Lower values indicate better spatial specificity, i.e., reduced off-target activation outside the biopsy needle region.

ficity is consistent across both benign and PCa cores and both centers, demonstrating that ANT not only boosts classification performance but also produces more spatially precise heatmap activations. Qualitative comparisons of these heatmap activations, demonstrating tighter spatial confinement with the proposed method, are shown in Figure 3.

5.5. Ablation Studies

5.5.1. ENCODER UPDATE STRATEGY

Table 7 reports the effect of varying which encoder blocks and parameters are updated during TTA. Our method updates all parameters of the first n blocks. We additionally




Figure 3: Representative heatmaps of model activations without TTA (PNF+) and with TTA (PNF+ + ANT) on the same biopsy cores. Activations represent predicted involvement, with higher values indicating greater likelihood of clinically significant cancer. Each example is annotated with grade group (GG) and corresponding involvement percentage.

Table 7: Ablation on encoder update strategy and auxiliary task. Mean AUC (%) across held-out test centers. Best in bold.

Update target	Position	Patient-Level (csPCa)	Core-Level (csPCa)	Core-Level (PCa)
<i>Layer position and parameter type</i>				
Full blocks	First n	82.1	84.4	72.7
Full blocks	Middle n	79.8	81.9	73.6
Full blocks	Last n	77.2	80.8	72.7
Full blocks	All blocks	74.9	79.6	70.2
Norm only	First n	76.4	80.9	72.8
Norm + bias	First n	80.4	82.2	73.0
Bias only	First n	78.8	80.6	71.8
<i>Auxiliary task signal (first n blocks, full update)</i>				
Reconstruction	First n	76.5	80.7	72.5
Segmentation (ours)	First n	82.1	84.4	72.7

compare against restricting updates to normalization parameters only or bias terms only, and against updating middle or late blocks.

Updating the first n encoder blocks yields the strongest performance across all evaluation metrics, outperforming updates applied to middle, late, or all blocks. Adapting the last n

Table 8: Ablation on number of TTA inner steps. Mean AUC (%) across held-out test centers. Best in bold.

Inner steps τ	Patient-Level (csPCa)	Core-Level (csPCa)	Core-Level (PCa)
0(w/o ANT)	78.5	81.5	72.7
2	76.0	81.1	72.6
5	78.9	83.4	72.5
10	82.1	84.4	72.7
20	78.2	78.0	72.0

blocks or the full network, consistently degrades performance, suggesting that these blocks capture task-specific representations that are sensitive to perturbation from the auxiliary objective. In contrast, early blocks appear to provide a more effective locus for adaptation.

Restricting updates to subsets of parameters within the first n blocks (e.g., normalization layers or bias terms) provides modest improvements but does not match the performance of full block updates. This indicates that partial parameter adaptation is insufficient, and that modifying the full set of parameters within early blocks is important for effective feature realignment. Finally, the results highlights a difference in performance based on the choice of auxiliary task, with anatomical alignment yielding stronger robustness to domain shift.

5.5.2. NUMBER OF TTA STEPS

Table 8 reports performance as a function of the number of inner optimization steps τ . Performance varies non-monotonically with the number of inner optimization steps τ . Increasing τ from 2 to 10 leads to consistent improvements across all metrics, with peak performance achieved at $\tau = 10$. However, further increasing the number of steps to 20 results in a noticeable decline, particularly at the core level. This suggests that a moderate number of adaptation steps is sufficient to achieve effective domain alignment, while excessive optimization may lead to overfitting or degradation of task-relevant features.

6. Discussion

This work shows that anatomically grounded test-time adaptation (TTA) via prostate segmentation improves both diagnostic performance and spatial specificity in micro-ultrasound prostate cancer detection under multi-center domain shift. Across unseen clinical centers and scanner generations, the proposed ANT method consistently outperforms both the non-adapted baseline and a strong state-of-the-art prostate cancer detection model (PNF+), while also improving anatomical alignment of model activations.

A consistent pattern emerges across all baseline TTA methods: performance improves at Center B (prevalence: 34.6%) but degrades at Center C (prevalence: 27.0%). This asymmetry is not straightforwardly explained by class prior shift alone as Center C has a prevalence closer to the source (16%) than Center B, yet shows greater degradation. Notably, ROID (Marsden et al., 2024), which explicitly incorporates prior correction to address class distribution shift, exhibits the same pattern, further suggesting that prevalence shift is not the primary driver. This observation may thus point to a more fundamental

limitation: adaptation signals that do not anchor to domain-invariant structure may remain susceptible to center-specific confounds, limiting their reliability under heterogeneous deployment conditions. The proposed ANT framework addresses this limitation by using prostate segmentation as a structured anatomical prior during test-time adaptation, which unlike entropy-based approaches, provides pixel-level spatial supervision that directly constrains adaptation to the organ of interest.

This interpretation is further supported by the reduction in activation leakage, defined as the ratio of average prostate heatmap activation to needle-region activation. As shown in Figure 2, while PNF+ exhibits a pronounced high-leakage tail indicating frequent off-target activations, PNF+ + ANT concentrates its distribution near zero across both benign and cancerous cores and both test centers. This suggests that ANT improves spatial specificity without sacrificing disease sensitivity as excessive off-target activation may reflect spurious correlations and undermine clinician trust. Additionally, the generalization of ANT to PNF+ without architecture-specific tuning further demonstrates that anatomical test-time adaptation is backbone-agnostic.

A central hypothesis of this work is that the prostate gland remains a consistently identifiable anatomical target across patient populations, clinical centers, and scanner generations, even as low-level image statistics vary substantially. The results of this study provide empirical support for this hypothesis: MicroSegNet generates pseudo-masks of sufficient quality across both test centers and the full range of patient populations encountered, as demonstrated by the pseudo-mask sensitivity analysis (Table 4), where ANT consistently improves over the unadapted baseline on csPCa AUC across all Dice tertiles including the low quality tertile (median Dice = 0.60). This suggests that the prostate gland is identifiable enough across the domain shifts encountered in this study to serve as a reliable anatomical supervision target, providing a more stable adaptation signal than entropy-based objectives.

A potential concern with ANT’s adaptation strategy is whether updating early encoder blocks while freezing downstream layers introduces a feature distribution mismatch that limits adaptation effectiveness. We attribute the consistent improvements observed across both test centers to two factors. First, gradients act on early blocks whose outputs propagate through many subsequent frozen layers before reaching the decoder, attenuating any representational mismatch at the decoder input. Second, the decoder was trained on augmented source domain images, providing implicit robustness to small feature distribution shifts introduced by test-time encoder updates. Together, these properties suggest that early-block adaptation via a segmentation objective produces feature shifts that are both gradual and within the decoder’s operating range, rather than introducing disruptive distributional discontinuities.

Limitations. ANT relies on a pretrained prostate segmentation model to generate pseudo-labels. While the pseudo-mask sensitivity analysis (Table 4) demonstrates that adaptation remains beneficial across the range of segmentation quality observed in this study, these results are limited to the domain shifts encountered between the source and two target centers. Whether the prostate gland remains sufficiently identifiable, and ANT therefore effective under more severe domain shifts, such as substantially different scanner generations, imaging protocols, or patient populations, remains an open question that future work should address.

References

- Hashim U Ahmed, Ahmed El-Shater Bosaily, Louise C Brown, Rhian Gabe, Richard Kaplan, Mahesh K Parmar, Yolanda Collaco-Moraes, Katie Ward, Richard G Hindley, Alex Freeman, et al. Diagnostic accuracy of multi-parametric mri and trus biopsy in prostate cancer (promis): a paired validating confirmatory study. *The Lancet*, 389(10071):815–822, 2017.
- M Jorge Cardoso, Wenqi Li, Richard Brown, Nic Ma, Eric Kerfoot, Yiheng Wang, Benjamin Murrey, Andriy Myronenko, Can Zhao, Dong Yang, et al. Monai: An open-source framework for deep learning in healthcare. *arXiv preprint arXiv:2211.02701*, 2022.
- Cheng Chen, Qi Dou, Hao Chen, Jing Qin, and Pheng Ann Heng. Unsupervised bidirectional cross-modality adaptation via deeply synergistic image and feature alignment for medical image segmentation. *IEEE transactions on medical imaging*, 39(7):2494–2505, 2020.
- Yurong Chen, Hui Zhang, Yaonan Wang, Weixing Peng, Wangdong Zhang, Q. M. Jonathan Wu, and Yimin Yang. D-bin: A generalized disentangling batch instance normalization for domain adaptation. *IEEE Transactions on Cybernetics*, 53(4):2151–2163, 2023. doi: 10.1109/TCYB.2021.3110128.
- Ziyang Chen, Yiwen Ye, Yongsheng Pan, and Yong Xia. Gradient alignment improves test-time adaptation for medical image segmentation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 2429–2437, 2025.
- Gregg Eure, Daryl Fanne, Jefferson Lin, Brian Wodlinger, and Sangeet Ghai. Comparison of conventional transrectal ultrasound, magnetic resonance imaging, and micro-ultrasound for visualizing prostate cancer in an active surveillance population: a feasibility study. *Canadian Urological Association Journal*, 13(3):E70, 2018.
- Chelsea Finn, Pieter Abbeel, and Sergey Levine. Model-agnostic meta-learning for fast adaptation of deep networks. In *International conference on machine learning*, pages 1126–1135. PMLR, 2017.
- Yaroslav Ganin, Evgeniya Ustinova, Hana Ajakan, Pascal Germain, Hugo Larochelle, François Laviolette, Mario March, and Victor Lempitsky. Domain-adversarial training of neural networks. *Journal of machine learning research*, 17(59):1–35, 2016.
- Mahdi Gilany, Paul Wilson, Andrea Perera-Ortega, Amoon Jamzad, Minh Nguyen Nhat To, Fahimeh Fooladgar, Brian Wodlinger, Purang Abolmaesumi, and Parvin Mousavi. Trusformer: improving prostate cancer detection from micro-ultrasound using attention and self-supervision. *International Journal of Computer Assisted Radiology and Surgery*, 18(7):1193–1200, 2023.
- Mahdi Gilany, Mohamed Harmanani, Paul Wilson, Minh Nguyen Nhat To, Amoon Jamzad, Fahimeh Fooladgar, Brian Wodlinger, Purang Abolmaesumi, and Parvin Mousavi. Calibrated diverse ensemble entropy minimization for robust test-time adaptation in prostate

- cancer detection. In *International Workshop on Machine Learning in Medical Imaging*, pages 361–371. Springer, 2024.
- Hao Guan and Mingxia Liu. Domain adaptation for medical image analysis: a survey. *IEEE Transactions on Biomedical Engineering*, 69(3):1173–1185, 2021.
- Ali Hatamizadeh, Yucheng Tang, Vishwesh Nath, Dong Yang, Andriy Myronenko, Bennett Landman, Holger R Roth, and Daguang Xu. Unetr: Transformers for 3d medical image segmentation. In *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pages 574–584, 2022.
- Hongxu Jiang, Muhammad Imran, Preethika Muralidharan, Anjali Patel, Jake Pensa, Muxuan Liang, Tarik Benidir, Joseph R Grajo, Jason P Joseph, Russell Terry, et al. Microseg-net: A deep learning approach for prostate segmentation on micro-ultrasound images. *Computerized Medical Imaging and Graphics*, 112:102326, 2024.
- Konstantinos Kamnitsas, Christian Baumgartner, Christian Ledig, Virginia Newcombe, Joanna Simpson, Andrew Kane, David Menon, Aditya Nori, Antonio Criminisi, Daniel Rueckert, et al. Unsupervised domain adaptation in brain lesion segmentation with adversarial networks. In *International conference on information processing in medical imaging*, pages 597–609. Springer, 2017.
- Adam Kinnaird, Ferdinand Luger, Hannes Cash, Sangeet Ghai, L Felipe Urdaneta-Salegui, Christian P Pavlovich, Joseph Brito, Neal D Shore, Julian P Struck, Martin Schostak, et al. Microultrasound-guided vs mri-guided biopsy for prostate cancer diagnosis: the optimum randomized clinical trial. *Jama*, 333(19):1679–1687, 2025.
- Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C Berg, Wan-Yen Lo, et al. Segment anything. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 4015–4026, 2023.
- Laurence Klotz, Giovanni Lughezzani, Davide Maffei, Andrea Sánchez, José Gregorio Pereira, Frédéric Staerman, Hannes Cash, Ferdinand Luger, Laurent Lopez, Rafael Sanchez-Salas, et al. Comparison of micro-ultrasound and multiparametric magnetic resonance imaging for prostate cancer: a multicenter, prospective analysis. *Canadian Urological Association Journal*, 15(1):E11, 2020.
- Da Li, Yongxin Yang, Yi-Zhe Song, and Timothy Hospedales. Learning to generalize: Meta-learning for domain generalization. In *Proceedings of the AAAI conference on artificial intelligence*, volume 32, 2018.
- Hao Li, Han Liu, Heinrich von Busch, Robert Grimm, Henkjan Huisman, Angela Tong, David Winkel, Tobias Penzkofer, Ivan Shabunin, Moon Hyung Choi, et al. Deep learning-based unsupervised domain adaptation via a unified model for prostate lesion detection using multisite biparametric mri datasets. *Radiology: Artificial Intelligence*, 6(5):e230521, 2024.

- Jian Liang, Ran He, and Tieniu Tan. A comprehensive survey on test-time adaptation under distribution shifts. *International Journal of Computer Vision*, 133(1):31–64, July 2024. ISSN 1573-1405. doi: 10.1007/s11263-024-02181-w. URL <http://dx.doi.org/10.1007/s11263-024-02181-w>.
- Jun Ma, Yuting He, Feifei Li, Lin Han, Chenyu You, and Bo Wang. Segment anything in medical images. *Nature Communications*, 15:654, 2024.
- Robert A Marsden, Mario Döbler, and Bin Yang. Universal test-time adaptation through weight ensembling, diversity weighting, and prior correction. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 2555–2565, 2024.
- Kinga Mate, Saleh Nedjim, Simon Bellucci, Cesar Boucault, Nael Ghaffar, Tracy Constantini, Fanni Marvanykovi, Pierre-Gilles Vestris, Yvanne Sadreux, Melanie Laguerre, et al. Prostate biopsy approach and complication rates. *Oncology Letters*, 26(3):375, 2023.
- Shuaicheng Niu, Jiaxiang Wu, Yifan Zhang, Yaofo Chen, Shijian Zheng, Peilin Zhao, and Mingkui Tan. Efficient test-time model adaptation without forgetting. In *The International Conference on Machine Learning*, 2022.
- Shuaicheng Niu, Jiaxiang Wu, Yifan Zhang, Zhiquan Wen, Yaofo Chen, Peilin Zhao, and Mingkui Tan. Towards stable test-time adaptation in dynamic wild world. In *International Conference on Learning Representations*, 2023.
- John Panzone, Timothy Byler, Gennady Bratslavsky, and Hanan Goldberg. Transrectal ultrasound in prostate cancer: current utilization, integration with mpMRI, hifu and other emerging applications. *Cancer Management and Research*, pages 1209–1228, 2022.
- Jake Pensa, Wayne Brisbane, Adam Kinnaird, David Kuppermann, Griffith Hughes, Derrick Ushko, Anthony Sisk, Leonard Marks, and Rory Geoghegan. Deep learning classification of prostate cancer on confidently labeled micro-ultrasound images. In *2024 46th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, pages 1–4. IEEE, 2024.
- Eduardo HP Pooch, Pedro Ballester, and Rodrigo C Barros. Can we trust deep learning based diagnosis? the impact of domain shift in chest radiograph classification. In *International workshop on thoracic image analysis*, pages 74–83. Springer, 2020.
- Joaquin Quiñero-Candela, Masashi Sugiyama, Anton Schwaighofer, and Neil D. Lawrence, editors. *Dataset Shift in Machine Learning*. MIT Press, Cambridge, MA, 2009.
- Prashanth Rawla. Epidemiology of prostate cancer. *World journal of oncology*, 10(2):63, 2019.
- Daniel Rohrbach, Brian Wodlinger, Jerrold Wen, Jonathan Mamou, and Ernest Feleppa. High-frequency quantitative ultrasound for imaging prostate cancer using a novel micro-ultrasound scanner. *Ultrasound in medicine & biology*, 44(7):1341–1354, 2018.

- Oriane Siméoni, Huy V Vo, Maximilian Seitzer, Federico Baldassarre, Maxime Oquab, Cijo Jose, Vasil Khalidov, Marc Szafraniec, Seungeun Yi, Michaël Ramamonjisoa, et al. Dinov3. *arXiv preprint arXiv:2508.10104*, 2025.
- Yu Sun, Xiaolong Wang, Zhuang Liu, John Miller, Alexei Efros, and Moritz Hardt. Test-time training with self-supervision for generalization under distribution shifts. In *International conference on machine learning*, pages 9229–9248. PMLR, 2020.
- Guillaume Vray, Devavrat Tomar, Xufeng Gao, Jean-Philippe Thiran, Evan Shelhamer, and Behzad Bozorgtabar. Reservoirtta: Prolonged test-time adaptation for evolving and recurring domains. *arXiv preprint arXiv:2505.14511*, 2025.
- Dequan Wang, Evan Shelhamer, Shaoteng Liu, Bruno Olshausen, and Trevor Darrell. Tent: Fully test-time adaptation by entropy minimization. *arXiv preprint arXiv:2006.10726*, 2020.
- Qin Wang, Olga Fink, Luc Van Gool, and Dengxin Dai. Continual test-time domain adaptation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 7201–7211, 2022.
- Jakob Wasserthal, Hanns-Christian Breit, Manfred T Meyer, Maurice Pradella, Daniel Hinck, Alexander W Sauter, Tobias Heye, Daniel T Boll, Joshy Cyriac, Shan Yang, et al. Totalsegmentator: robust segmentation of 104 anatomic structures in ct images. *Radiology: Artificial Intelligence*, 5(5):e230024, 2023.
- Emma Willis, Tarek Elghareb, Paul F. R. Wilson, Minh Nguyen Nhat To, Mohammad Mahdi Abootorabi, Amoon Jamzad, Brian Wodlinger, Parvin Mousavi, and Purang Abolmaesumi. Guide-us: Grade-informed unpaired distillation of encoder knowledge from histopathology to micro-ultrasound, 2026. URL <https://arxiv.org/abs/2602.19005>.
- Paul FR Wilson, Minh Nguyen Nhat To, Amoon Jamzad, Mahdi Gilany, Mohamed Harmanani, Tarek Elghareb, Fahimeh Fooladgar, Brian Wodlinger, Purang Abolmaesumi, and Parvin Mousavi. Prostnfound: integrating foundation models with ultrasound domain knowledge and clinical context for robust prostate cancer detection. In *International conference on medical image computing and computer-assisted intervention*, pages 499–509. Springer, 2024.
- Paul FR Wilson, Mohamed Harmanani, Minh Nguyen Nhat To, Amoon Jamzad, Tarek Elghareb, Zhuoxin Guo, Adam Kinnaird, Brian Wodlinger, Purang Abolmaesumi, and Parvin Mousavi. Prostnfound+: A prospective study using medical foundation models for prostate cancer detection. *International Journal of Computer Assisted Radiology and Surgery*, pages 1–10, 2026.
- Jianghao Wu, Xiangde Luo, Yubo Zhou, Lianming Wu, Guotai Wang, and Shaoting Zhang. A3-tta: Adaptive anchor alignment test-time adaptation for image segmentation. *IEEE Transactions on Image Processing*, 34:8511–8522, 2025.
- Jee Seok Yoon, Kwanseok Oh, Yooseung Shin, Maciej A Mazurowski, and Heung-Il Suk. Domain generalization for medical image analysis: A review. *Proceedings of the IEEE*, 112(10):1583–1609, 2024.

- Wenjing Yu, Shuo Jiang, Yifei Chen, Shuo Chang, Yuanhan Wang, Beining Wu, Jie Dong, Mingxuan Liu, Shenghao Zhu, Feiwei Qin, et al. A large scale benchmark for test time adaptation methods in medical image segmentation. *arXiv preprint arXiv:2512.02497*, 2025.
- Ling Zhang, Xiaosong Wang, Dong Yang, Thomas Sanford, Stephanie Harmon, Baris Turkbey, Bradford J Wood, Holger Roth, Andriy Myronenko, Daguang Xu, et al. Generalizing deep learning for medical image segmentation to unseen domains via deep stacked transformation. *IEEE transactions on medical imaging*, 39(7):2531–2540, 2020.
- Marvin Zhang, Sergey Levine, and Chelsea Finn. Memo: Test time robustness via adaptation and augmentation. *Advances in neural information processing systems*, 35:38629–38642, 2022.
- Kaiyang Zhou, Yongxin Yang, Yu Qiao, and Tao Xiang. Mixstyle neural networks for domain generalization and adaptation. *International Journal of Computer Vision*, 132(3):822–836, 2024.

Appendix A. Implementation Details

A.1. Cancer Detection Model

A.1.1. TRAINING CONFIGURATION

All experiments were run using PyTorch 2.1 on a single NVIDIA L40S GPU (48 GB). The detection model was trained for 35 epochs with a batch size of 4 using the AdamW optimizer ($\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$, weight decay = 0.01), with a fixed random seed of 30. The encoder (DINOv3 ViT-L/16) was initialized with self-supervised pretrained weights and was fine-tuned with a learning rate of 10^{-5} , and the UNETR decoder with a learning rate of 10^{-4} , both with cosine annealing scheduling and a 5-epoch linear warmup. These hyperparameters follow those reported in ProstNFound+ (Wilson et al., 2026).

A.1.2. DATA AUGMENTATION

Augmentation was applied to the training set only. Two transformations were applied sequentially: random affine translation with a maximum displacement of 20% in each spatial dimension, and random resized cropping to 512×512 pixels with scale drawn uniformly from (0.3, 1.0), applied with probability 0.5. The needle mask and prostate mask were transformed consistently with the image to preserve spatial correspondence. Images were normalized to $[0, 1]$ by min-max scaling and then standardized with mean (0, 0, 0) and standard deviation (1, 1, 1).

A.1.3. CLASS AND MASK DECODER

Following the UNETR encoder-decoder, the spatial feature representations are passed to a two-way transformer module (Kirillov et al., 2023) that performs bidirectional cross-attention between image embeddings and learnable query tokens. This produces contextualized representations that are then routed to two task-specific decoders.

A.1.4. MASK DECODER

The mask decoder receives the two-way transformer output and produces a dense spatial heatmap $P \in [0, 1]^{H \times W}$ over the image. It consists of a sequence of transposed convolutional upsampling blocks that progressively increase spatial resolution from the compressed transformer output back to the input image dimensions. The heatmap is activated with sigmoid to produce per-pixel cancer probability estimates. Core-level cancer scores are derived by averaging heatmap activations within the needle trace region.

A.1.5. CLASS DECODER

The class decoder receives the same two-way transformer output and produces an image-level classification score for clinically significant PCa. A learnable classification token attends over the image embeddings via cross-attention, and the resulting token representation is projected through a small MLP to produce a scalar logit $z_{\text{cls}} \in \mathbb{R}$, which is activated with sigmoid to give the predicted csPCa probability $\hat{y}_{\text{cls}} \in [0, 1]$.

The two decoders are trained jointly via the multi-task loss described in Section 3.4.2, with the mask decoder supervised by core-level involvement labels and the class decoder

supervised by binary csPCa labels. At inference, the class decoder score is used for patient-level and image-level evaluation, while the mask decoder heatmap is used for core-level evaluation via needle trace averaging. During TTA, both decoders remain frozen.

A.1.6. EVALUATION METRICS

Performance was evaluated using AUC of the receiver operating characteristic curve at three levels: patient-level csPCa, core-level csPCa ($\text{GG} \geq 3$ vs. all other cores), and core-level PCa (any cancer vs. benign). Core-level scores are derived by averaging heatmap activations along the biopsy needle trace mask. For patient-level evaluation, the patient-level prediction is the mean of the model’s core-level scores across all cores for that patient. Patient-level csPCa AUC reflects the classification of patients designated as csPCa ($\text{GG} \geq 3$) versus non-csPCa based on pathological grading.

A.2. Test-Time Adaptation

A.2.1. SEGMENTATION HEAD ARCHITECTURE

The segmentation head attached to the encoder for TTA consists of two 3×3 convolutional layers each followed by GroupNorm (32 groups) and ReLU activation, and a final 1×1 projection to a single output channel:

$$p_\psi : \mathbb{R}^{C \times H' \times W'} \rightarrow \mathbb{R}^{1 \times H' \times W'} \quad (8)$$

where $C = 1024$ is the encoder feature dimension and the intermediate hidden dimension is 256. The head is randomly initialized at the start of the test set and is never reset throughout inference.

A.2.2. TTA OPTIMIZER

The encoder adaptation and segmentation head use separate Adam optimizers. The encoder optimizer uses a learning rate of $\eta_{\text{enc}} = \eta_{\text{TTT}} \times 0.1$ and the segmentation head optimizer uses $\eta_{\text{head}} = \eta_{\text{TTT}}$. Both optimizers are stateful and persist across all test samples, accumulating first and second moment estimates throughout the test set. Test samples follow a fixed sequential order by patient ID. For the sensitivity analysis results reported in (Table 11), the following random seeds were used: 10, 20, 40, 60, and 70.

A.2.3. TTA LOSS

The TTA objective is a combined Dice and binary cross-entropy loss applied between the predicted segmentation $\hat{s}_j$ and the pseudo-mask $\hat{m}_j$:

$$\mathcal{L}_{\text{TTA}} = \mathcal{L}_{\text{Dice}}(\hat{s}_j, \hat{m}_j) + \lambda \mathcal{L}_{\text{BCE}}(\hat{s}_j, \hat{m}_j), \quad (9)$$

Segmentation quality during TTA is monitored using the Dice similarity coefficient between binarized predictions ($\hat{s}_j > 0.5$) and binarized pseudo-masks ($\hat{m}_j > 0.5$), computed using the MONAI `DiceMetric` (Cardoso et al., 2022) with background included. In our setting, we fix $\lambda = 1$.

Table 9: Selected TTA hyperparameters.

Hyperparameter	Value
TTA learning rate (η_{TTT})	10^{-4}
Encoder learning rate (η_{enc})	10^{-5}
Inner steps (τ)	10
Blocks to adapt (n , DINOv3)	6
Blocks to adapt (n , MedSAM)	3
Segmentation head hidden dim	256

A.2.4. BLOCK SELECTION

The encoder blocks are partitioned into three equal groups (first third, middle third, last third) based on total block count L :

$$\text{first} = \text{blocks}[0 : L/3], \quad \text{middle} = \text{blocks}[L/3 : 2L/3], \quad \text{last} = \text{blocks}[2L/3 : L] \quad (10)$$

For the DINOv3 ViT-L/16 encoder ($L = 24$), the first third comprises blocks 0 – 7. The optimal number of blocks to adapt was $n = 6$, corresponding to the first 6 blocks (25% of the encoder). For the MedSAM encoder in ProstNFound+ ($L = 12$), the equivalent proportion yields $n = 3$ blocks. When restricting to normalization parameters only, only the affine parameters (γ , β) of LayerNorm, BatchNorm, and GroupNorm layers are included. When restricting to bias terms only, only the bias parameters of Linear and Conv2d layers are included.

A.2.5. HYPERPARAMETER SELECTION

TTA hyperparameters were selected on a held-out split of the source domain training data (80% train / 20% validation, patient-level split across 693 patients). The selection criterion was core-level PCa AUC on the validation split. The selected hyperparameters were fixed for all LOCO evaluation folds without center-specific tuning. Optimal values are summarized in Table 9.

A.3. Real-Time Inference Latency

Table 10 reports wall-clock inference time per image for all methods on a single NVIDIA L40S GPU (48 GB). ANT requires 0.90 ± 0.06 seconds per image ($\tau = 10$ inner optimisation steps), compared to 0.07 ± 0.01 seconds for standard inference without adaptation. While ANT is slower than all baseline methods, its sub-second inference time is compatible with the procedural cadence of micro-ultrasound guided biopsy, where needle positioning and firing inherently require several seconds per core.

A.4. Long-Term Sequential Stability

To assess ordering sensitivity, ANT was evaluated under five independent random permutations of the test set and compared against five runs under the fixed original ordering,

Table 10: Mean wall-clock inference time per image (seconds $\pm$ std) on a single NVIDIA L40S GPU (48 GB). ANT includes $\tau = 10$ inner optimisation steps per image.

Method	Time (s)
No adaptation	0.07 ± 0.01
TENT	0.11 ± 0.03
EATA	0.13 ± 0.01
ROID	0.19 ± 0.03
SAR	0.22 ± 0.02
CoTTA	0.26 ± 0.02
MEMO	0.29 ± 0.02
ANT (ours)	0.90 ± 0.06

Table 11: Ordering sensitivity analysis. Mean AUC (%) $\pm$ std across five random seeds for fixed and randomly permuted test set orderings.

Ordering	Patient-Level csPCa AUC (%)		Core-Level csPCa AUC (%)		Core-Level PCa AUC (%)	
	B	C	B	C	B	C
Fixed order	79.9 ± 0.03	84.1 ± 0.08	81.1 ± 0.02	88.0 ± 0.06	73.1 ± 0.03	72.3 ± 0.03
Random order	80.0 ± 0.04	82.3 ± 0.03	80.6 ± 0.04	85.8 ± 0.08	72.9 ± 0.01	72.0 ± 0.05

Table 12: Calibration analysis on Center B ($n = 1,118$ cores). Expected Calibration Error (ECE) is reported for the unadapted baseline (PNF+) and ANT. $\Delta\text{ECE} = \text{ANT ECE} - \text{Baseline ECE}$; negative values indicate improved calibration.

Subgroup (Involvement)	n	Baseline ECE (%)	ANT ECE (%)	ΔECE (%)
Overall	1,118	7.6	2.7	-4.9
Low	161	29.7	32.0	+2.2
Mid	93	19.6	16.5	-3.0
High	131	14.3	13.8	-0.5

reporting mean AUC and standard deviation across seeds (Table 11). Performance is stable across orderings, with differences of less than 2 AUC points across all metrics and centres, and overlapping standard deviations in all cases, demonstrating that ANT is robust to the order in which test images are presented during episodic adaptation.

A.5. Calibration Analysis

We assessed the effect of ANT on model calibration using Expected Calibration Error (ECE) on Center B, which provides the most representative subgroup distribution. ANT substantially improves overall calibration relative to the PNF+ baseline (ECE: 2.7% vs. 7.6%, $\Delta\text{ECE} = -4.9\%$). For cancer-positive cores stratified by involvement tertile, ANT

maintains or improves calibration in the mid ($\Delta\text{ECE} = -3.0\%$) and high ($\Delta\text{ECE} = -0.5\%$) involvement subgroups. A modest increase in ECE is observed in the low involvement subgroup ($\Delta\text{ECE} = +2.2\%$), suggesting slight overconfidence on cores with minimal cancer extent, which is consistent with the inherent difficulty of detecting low-involvement cancer under domain shift.